

FinalProject_group78

June 12, 2019

1 COGS 108 Final Project

2 Overview

For our project, we decided to see if we could predict the quality of wine given its ingredients. Since wine is so easily and heavily influenced by a number of factors that we can't control for (i.e humidity, grape quality, etc), we wanted to focus on its internal workings. We hypothesized that the volume of certain features, namely alcohol percentage, volatile acidity and citric acid, were crucial in determining a wine's quality rating. Using two datasets to test our prediction, we found that these, as well as the other characteristics, had no individual effect on overall rating, but rather, the rating was derived from a combination of different features.

3 Group Members - Names and ID

- Alysha Ali (A15345048)
- Tsai-Ying Ying Chuang (A15215376)
- Zheng Zeng (A14679117)
- Melody Yu (A14599481)
- Ai-Ting Hsieh (A13595395)

4 Part 1 - Introduction & Background

Since its creation 6000 years ago, wine has been adapted widely — from an instrument for religious ceremonies to a bargaining tool between countries (<https://vinepair.com/wine-colonized-world-wine-history/#5>). Despite the variety in usage, wine has constantly been regarded as a social device and has only risen in popularity (<https://www.svb.com/globalassets/library/images/svb-2018-wine-report.pdf>). That's why, with the combination of the quickly-approaching legal drinking age for several of our group members and its increasing appearance in social media(<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4432862/>), we thought it was important to determine how the quality of a wine is determined. Specifically, we aim to discover exactly what ingredients are crucial in deriving a wine's quality rating.

We came to our hypothesis by first researching wine characteristics (<https://winefolly.com/review/wine-characteristics/>). From there, we saw that certain elements, like citric acid content and volatile acidity, directly influenced a wine's sweetness. From there, we looked at ranges and eliminated factors like water density, which only ranged from 0.99

to 1.00 and therefore could not have as large an impact on quality as other factors, like total sulfur dioxide for example, which had a range of over 200.

This information is important for all consumers, from pre-teens getting their first sip of wine to experienced connoisseurs and businesses. With this data analysis, instead of spending hours trying to choose a wine everyone will love, party hosts and other entertainers will know what the most favored wine features are and be able to buy a delicious spirit in a snap. For instance, knowing what chemicals make a wine popular will reduce how much time winemakers spend in perfecting their formula and ultimately maximize efficiency and increase profit in the long run. Thus, this information can affect more than just consumers.

Since these datasets were previously the subject of a competition hosted by the University of California, Irvine, there have been several other projects that have analyzed similar questions and/or have used the same datasets. Most notably, is another study that also attempted to predict the quality of wine (<https://www.kaggle.com/vishalyo990/prediction-of-quality-of-wine>). However, they differed in their methods used and the overall accuracy of their machine learning model.

4.1 Our question

Can we predict the quality of a wine given a list of its physicochemical features?

4.2 What we think

We hypothesize that we can predict the quality of the wine by its individual features. Especially, we predict that the lower volatile acidity, medium alcohol percentage, and higher citric acid will create the best quality of the wine.

5 Part 2 - Data Wrangling

To answer our question, we will be using two datasets: one on red wines and one on white wines (source: <https://archive.ics.uci.edu/ml/datasets/wine+quality>). The red wine dataset had 1601 observations, while the white wine dataset had 4900. These datasets are identical in the features they analyze, namely a list of physiochemical ingredients in wine:

- fixed acidity - most acids involved with wine or fixed or nonvolatile (do not evaporate readily)
- volatile acidity - the amount of acetic acid in wine, which at too high of levels can lead to an unpleasant, vinegar taste
- citric acid - found in small quantities, citric acid can add 'freshness' and flavor to wines
- residual sugar - the amount of sugar remaining after fermentation stops, it's rare to find wines with less than 1 gram/liter and wines with greater than 45 grams/liter are considered sweet
- chlorides - the amount of salt in the wine
- free sulfur dioxide - the free form of SO₂ exists in equilibrium between molecular SO₂ (as a dissolved gas) and bisulfite ion; it prevents microbial growth and the oxidation of wine

- total sulfur dioxide - amount of free and bound forms of SO₂; in low concentrations, SO₂ is mostly undetectable in wine, but at free SO₂ concentrations over 50 ppm, SO₂ becomes evident in the nose and taste of wine
- density - the density of wine is close to that of water depending on the percent alcohol and sugar content
- pH - describes how acidic or basic a wine is on a scale from 0 (very acidic) to 14 (very basic); most wines are between 3-4 on the pH scale
- sulphates - a wine additive which can contribute to sulfur dioxide gas (SO₂) levels, which acts as an antimicrobial and antioxidant
- alcohol - the percent alcohol content of the wine
- quality - output variable (based on sensory data, score between 0 and 10)

We combined the two in several ways during our process. For instance, in one part, we began by graphing the individual features together in order to see a comparison between white and red wines. In another part, we used random samples from both datasets in order to train and test our algorithm.

5.1 Setup

```
In [4]: import pandas as pd
import numpy as np
import pylab
import requests
import bs4
from bs4 import BeautifulSoup
import matplotlib.pyplot as plt
from sklearn.neighbors import KDTree
from sklearn.utils import shuffle
from pandas.plotting import scatter_matrix

In [5]: #Here we import the two datasets that we will be using in our final project
df_red = pd.read_csv('wineQualityReds.csv')
df_white = pd.read_csv('wineQualityWhites.csv')

# print out the column names
print(df_red.columns.values)

['Unnamed: 0' 'fixed.acidity' 'volatile.acidity' 'citric.acid'
'residual.sugar' 'chlorides' 'free.sulfur.dioxide' 'total.sulfur.dioxide'
'density' 'pH' 'sulphates' 'alcohol' 'quality']
```

6 Part 3 - Data Cleaning, Processing & Visualization

Here, we will clean up our data before we start analyzing. We will remove rows with missing data, and delete columns that are not useful. More explanation will be included as we clean up the data.

```
In [6]: df_red[:5]
```

```
Out[6]:
```

	Unnamed: 0	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	\
0	1	7.4	0.70	0.00	1.9	
1	2	7.8	0.88	0.00	2.6	
2	3	7.8	0.76	0.04	2.3	
3	4	11.2	0.28	0.56	1.9	
4	5	7.4	0.70	0.00	1.9	

	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	\
0	0.076	11.0	34.0	0.9978	3.51	
1	0.098	25.0	67.0	0.9968	3.20	
2	0.092	15.0	54.0	0.9970	3.26	
3	0.075	17.0	60.0	0.9980	3.16	
4	0.076	11.0	34.0	0.9978	3.51	

	sulphates	alcohol	quality
0	0.56	9.4	5
1	0.68	9.8	5
2	0.65	9.8	5
3	0.58	9.8	6
4	0.56	9.4	5

```
In [7]: df_red.shape
```

```
Out[7]: (1599, 13)
```

Drop any row with missing values

```
In [8]: df_red = df_red.dropna()
```

Drop the column labeled "Unnamed: 0"

```
In [9]: df_red = df_red.drop(['Unnamed: 0'], axis=1)
```

```
In [10]: df_red.shape
```

```
Out[10]: (1599, 12)
```

```
In [11]: df_red[:5]
```

```
Out[11]:
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	

2	7.8	0.76	0.04	2.3	0.092
3	11.2	0.28	0.56	1.9	0.075
4	7.4	0.70	0.00	1.9	0.076

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates \
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

In [12]: df_white[:5]

```
Out[12]: Unnamed: 0  fixed.acidity  volatile.acidity  citric.acid  residual.sugar \
0          1          7.0          0.27          0.36          20.7
1          2          6.3          0.30          0.34          1.6
2          3          8.1          0.28          0.40          6.9
3          4          7.2          0.23          0.32          8.5
4          5          7.2          0.23          0.32          8.5
```

	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density	pH \
0	0.045	45.0	170.0	1.0010	3.00
1	0.049	14.0	132.0	0.9940	3.30
2	0.050	30.0	97.0	0.9951	3.26
3	0.058	47.0	186.0	0.9956	3.19
4	0.058	47.0	186.0	0.9956	3.19

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

In [13]: df_white.shape

Out[13]: (4898, 13)

Drop any rows with missing values

In [14]: df_white = df_white.dropna()

Drop the column labeled as "Unnamed: 0"

```
In [15]: df_white = df_white.drop(['Unnamed: 0'], axis=1)
```

```
In [16]: df_white.shape
```

```
Out[16]: (4898, 12)
```

```
In [17]: df_white[:5]
```

```
Out[17]:
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	\
0	7.0	0.27	0.36	20.7	0.045	
1	6.3	0.30	0.34	1.6	0.049	
2	8.1	0.28	0.40	6.9	0.050	
3	7.2	0.23	0.32	8.5	0.058	
4	7.2	0.23	0.32	8.5	0.058	

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates	\
0	45.0	170.0	1.0010	3.00	0.45	
1	14.0	132.0	0.9940	3.30	0.49	
2	30.0	97.0	0.9951	3.26	0.44	
3	47.0	186.0	0.9956	3.19	0.40	
4	47.0	186.0	0.9956	3.19	0.40	

	alcohol	quality
0	8.8	6
1	9.5	6
2	10.1	6
3	9.9	6
4	9.9	6

We want to predict wine quality based on their ingredients therefore the information we need for each wine instance is the list of ingredient values and the quality. In the above cells, we first dropped any rows with missing values to ensure that we are using a dataset that contains exactly what we need. Since this is already a very clean dataset, none of the rows were dropped. The column labeled “Unnamed: 0” does not seem to be useful at all since we also have the index of each row. Therefore, we dropped the entire column. Now, the resulting dataframe only contains information we need.

Since the white wine dataset contains more instances, we will only take 1599 instances from it randomly to match the number of instances of red wine.

```
In [18]: df_white = shuffle(df_white).reset_index(drop = True)
```

```
df_white = df_white[:1599]
```

Here we will combine two wine sets.

```
In [19]: # appending datasets of both white wine and red wine
df_combined = df_red.append(df_white, ignore_index = True)
df_combined = shuffle(df_combined)
df_combined = df_combined.reset_index(drop=True)

df_combined.shape
```

Out [19]: (3198, 12)

```
In [20]: desc = df_combined.describe()
desc
```

```
Out [20]:
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	\
count	3198.000000	3198.000000	3198.000000	3198.000000	
mean	7.593340	0.403094	0.302489	4.489290	
std	1.549998	0.192557	0.164095	4.192491	
min	4.200000	0.080000	0.000000	0.600000	
25%	6.600000	0.250000	0.210000	1.900000	
50%	7.200000	0.360000	0.300000	2.400000	
75%	8.200000	0.530000	0.400000	6.000000	
max	15.900000	1.580000	1.000000	31.600000	

	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density	\
count	3198.000000	3198.000000	3198.000000	3198.000000	
mean	0.066540	25.200281	91.969669	0.995411	
std	0.042143	16.418179	59.627286	0.002819	
min	0.012000	1.000000	6.000000	0.987420	
25%	0.043000	12.000000	37.000000	0.993443	
50%	0.059000	22.000000	88.000000	0.995800	
75%	0.080000	35.000000	136.000000	0.997400	
max	0.611000	112.000000	366.500000	1.010300	

	pH	sulphates	alcohol	quality
count	3198.000000	3198.000000	3198.000000	3198.000000
mean	3.249912	0.572745	10.461024	5.761101
std	0.163454	0.167567	1.151039	0.851406
min	2.720000	0.230000	8.000000	3.000000
25%	3.140000	0.460000	9.500000	5.000000
50%	3.240000	0.550000	10.200000	6.000000
75%	3.360000	0.650000	11.200000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
In [21]: corrs = df_combined.corr()
corrs
```

```
Out [21]:
```

	fixed.acidity	volatile.acidity	citric.acid	\
fixed.acidity	1.000000	0.171727	0.407795	
volatile.acidity	0.171727	1.000000	-0.458312	
citric.acid	0.407795	-0.458312	1.000000	
residual.sugar	-0.164296	-0.281011	0.164973	
chlorides	0.296037	0.361918	0.059043	
free.sulfur.dioxide	-0.334627	-0.397302	0.132845	
total.sulfur.dioxide	-0.372436	-0.454174	0.197656	
density	0.556108	0.324357	0.111073	
pH	-0.292510	0.331062	-0.424707	
sulphates	0.345303	0.200511	0.113168	

alcohol	-0.090773	-0.095184	0.034802
quality	-0.043420	-0.333890	0.158044

	residual.sugar	chlorides	free.sulfur.dioxide	\
fixed.acidity	-0.164296	0.296037	-0.334627	
volatile.acidity	-0.281011	0.361918	-0.397302	
citric.acid	0.164973	0.059043	0.132845	
residual.sugar	1.000000	-0.195536	0.489241	
chlorides	-0.195536	1.000000	-0.257042	
free.sulfur.dioxide	0.489241	-0.257042	1.000000	
total.sulfur.dioxide	0.555424	-0.326883	0.768073	
density	0.349580	0.376899	-0.092566	
pH	-0.293666	0.018465	-0.180862	
sulphates	-0.236341	0.462941	-0.243341	
alcohol	-0.280145	-0.227448	-0.152875	
quality	0.023986	-0.193329	0.090369	


	total.sulfur.dioxide	density	pH	sulphates	\
fixed.acidity	-0.372436	0.556108	-0.292510	0.345303	
volatile.acidity	-0.454174	0.324357	0.331062	0.200511	
citric.acid	0.197656	0.111073	-0.424707	0.113168	
residual.sugar	0.555424	0.349580	-0.293666	-0.236341	
chlorides	-0.326883	0.376899	0.018465	0.462941	
free.sulfur.dioxide	0.768073	-0.092566	-0.180862	-0.243341	
total.sulfur.dioxide	1.000000	-0.136670	-0.296915	-0.333072	
density	-0.136670	1.000000	0.022438	0.326821	
pH	-0.296915	0.022438	1.000000	0.139735	
sulphates	-0.333072	0.326821	0.139735	1.000000	
alcohol	-0.202789	-0.618483	0.144421	0.017505	
quality	0.007226	-0.282348	-0.014176	0.066996	

	alcohol	quality
fixed.acidity	-0.090773	-0.043420
volatile.acidity	-0.095184	-0.333890
citric.acid	0.034802	0.158044
residual.sugar	-0.280145	0.023986
chlorides	-0.227448	-0.193329
free.sulfur.dioxide	-0.152875	0.090369
total.sulfur.dioxide	-0.202789	0.007226
density	-0.618483	-0.282348
pH	0.144421	-0.014176
sulphates	0.017505	0.066996
alcohol	1.000000	0.439750
quality	0.439750	1.000000

By reading the above correlation table, it seems like none of the feature correlates with the wine's quality, we will plot out some matrices to see if that really is the case.

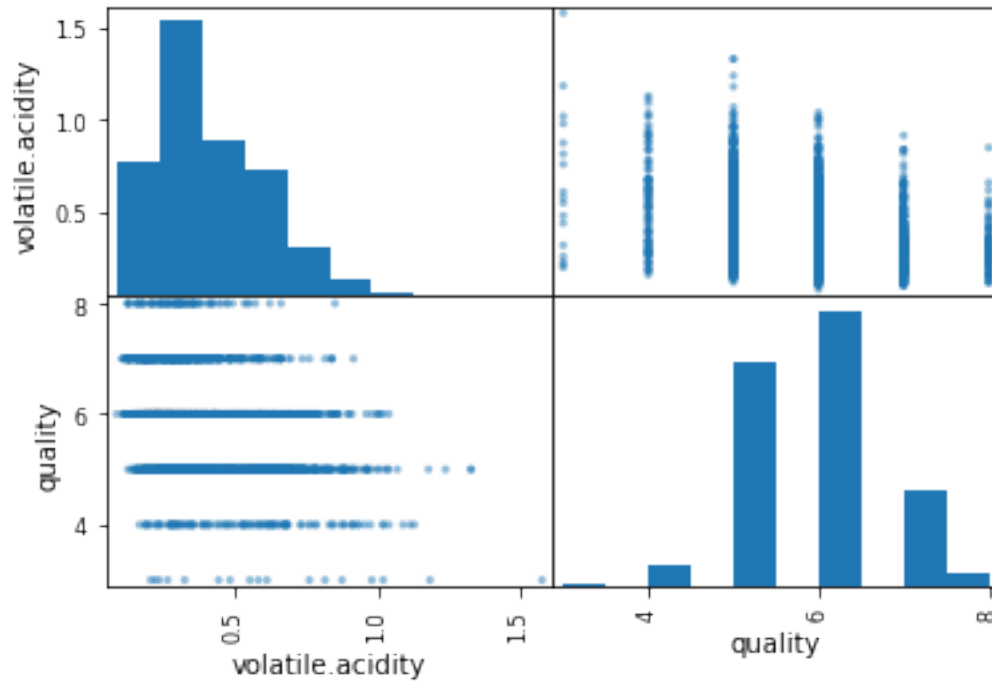
In [22]: `scatter_matrix(df_combined[['fixed.acidity','quality']])`


```
f4 = plt.gcf()
```



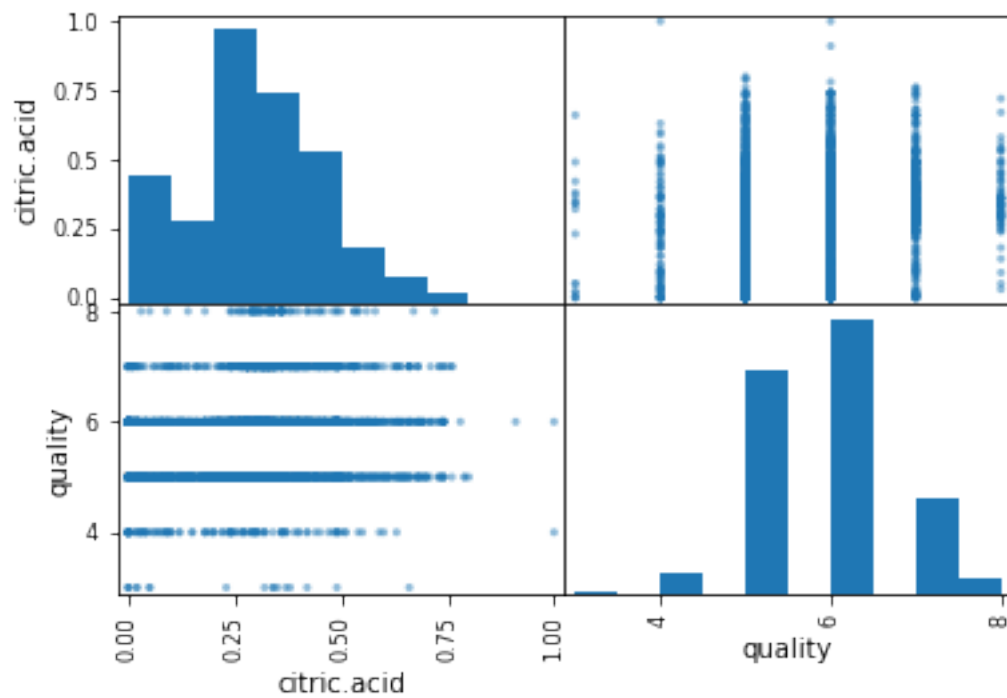
```
In [23]: scatter_matrix(df_combined[['volatile.acidity','quality']])
```

```
f4 = plt.gcf()
```



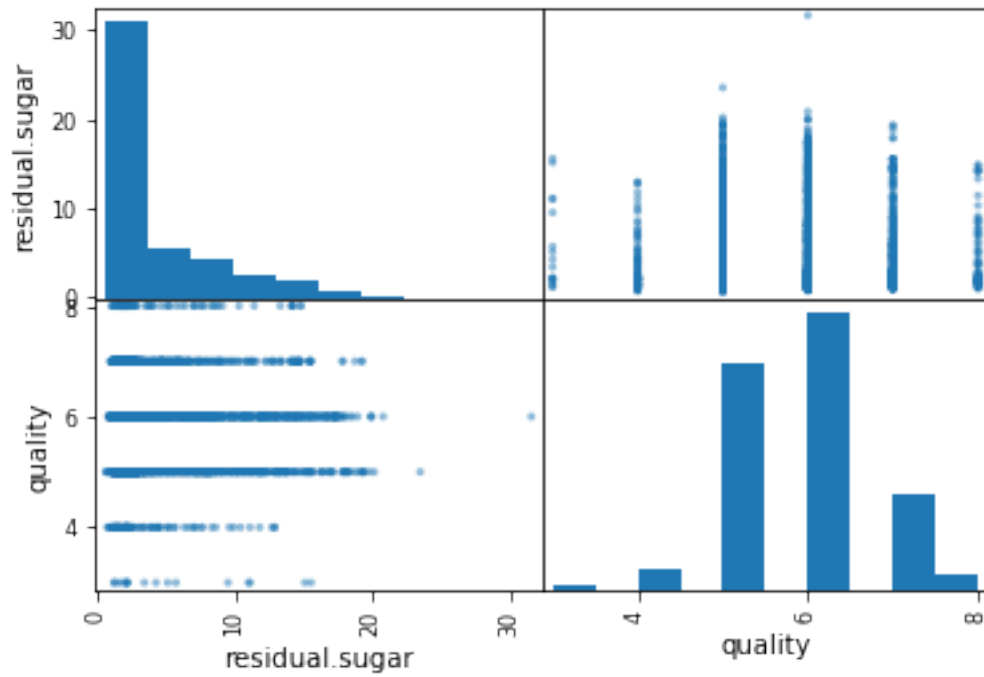
```
In [24]: scatter_matrix(df_combined[['citric.acid', 'quality']])
```

```
f4 = plt.gcf()
```



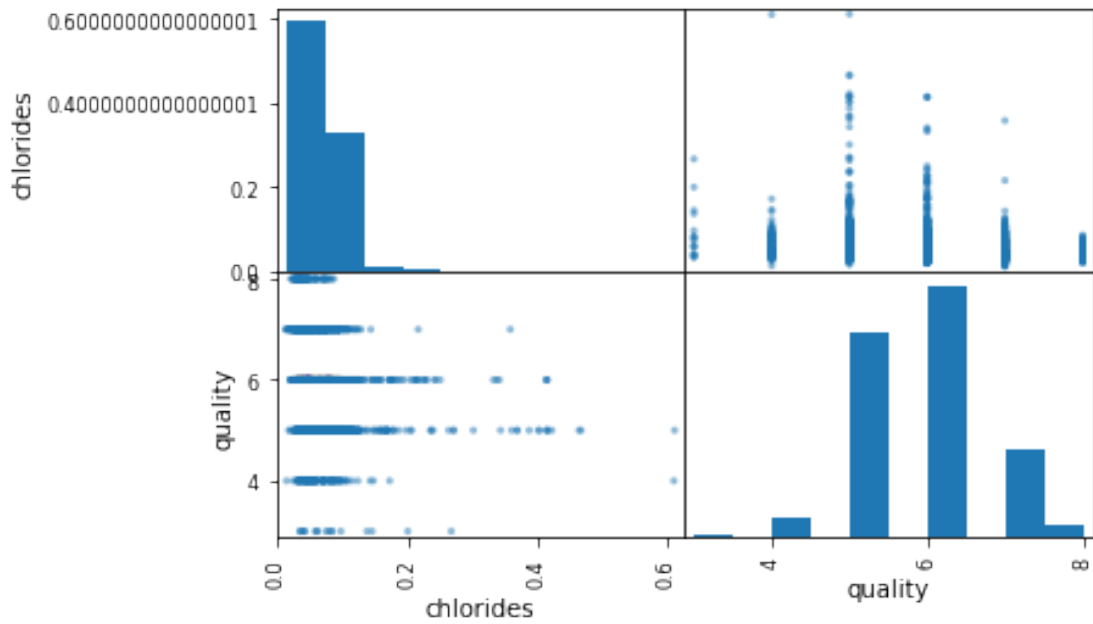
```
In [25]: scatter_matrix(df_combined[['residual.sugar','quality']])
```

```
f4 = plt.gcf()
```




```
In [26]: scatter_matrix(df_combined[['chlorides','quality']])
```

```
f4 = plt.gcf()
```



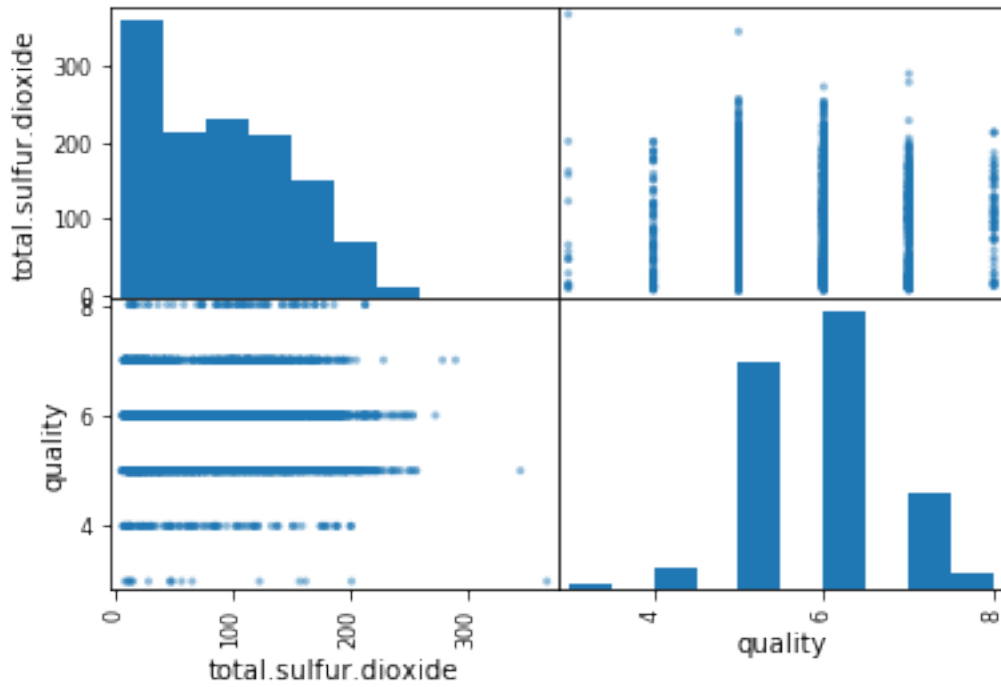
```
In [27]: scatter_matrix(df_combined[['free.sulfur.dioxide','quality']])
```

```
f4 = plt.gcf()
```




```
In [28]: scatter_matrix(df_combined[['total.sulfur.dioxide','quality']])

f4 = plt.gcf()
```




```
In [29]: scatter_matrix(df_combined[['density','quality']])

f4 = plt.gcf()
```




```
In [30]: scatter_matrix(df_combined[['pH', 'quality']])
```

```
f4 = plt.gcf()
```




```
In [31]: scatter_matrix(df_combined[['sulphates', 'quality']])
```

```
f4 = plt.gcf()
```



```
In [32]: scatter_matrix(df_combined[['alcohol', 'quality']])
```

```
f4 = plt.gcf()
```



Since none of the feature seems to be highly correlated to the quality, we want to combine them all and see if doing so would allow us to predict the wine rating more accurately

```
In [33]: df_combined.shape
```

```
Out[33]: (3198, 12)
```

Similar to A5, we will split our dataset into training set (80% of total) and testing set (20% of total).

```
In [34]: num_training = int(3198*.8)
         num_training
```

```
Out[34]: 2558
```

```
In [35]: df_training = df_combined[:num_training]
         df_testing = df_combined[num_training:]
```

The following show what our training set looks like if we ignore the kind of the wine.

```
In [36]: df_training[:5]
```

```
Out[36]:
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	\
0	6.5	0.28	0.33	15.7	0.053	
1	8.3	0.31	0.39	2.4	0.078	
2	6.8	0.18	0.37	1.6	0.055	
3	6.4	0.30	0.33	5.2	0.050	

4	9.1	0.28	0.48	1.8	0.067
---	-----	------	------	-----	-------

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates \
0	51.0	190.0	0.99780	3.22	0.51
1	17.0	43.0	0.99444	3.31	0.77
2	47.0	154.0	0.99340	3.08	0.45
3	30.0	137.0	0.99304	3.26	0.58
4	26.0	46.0	0.99670	3.32	1.04

	alcohol	quality
0	9.7	6
1	12.5	7
2	9.1	5
3	11.1	5
4	10.6	6

Now, we will include some functions that can be useful later.

```
In [37]: #takes in a dataframe and a list of list of indices,
#and return the predicted rate(quality) for each whine in the dataframe
#based on the list of corresponding indices
def predict_rate(df, inds):
    rate_predict = []
    for i in inds:
        rateMap = {}
        rate_mode = 0
        for j in i:
            curr = df.loc[j, 'quality']
            if curr in rateMap.keys():
                rateMap[curr]+=1
            else:
                rateMap[curr] = 1
            if rateMap[curr] > rate_mode:
                rate_mode = curr
        rate_predict.append(rate_mode)
    return rate_predict

In [38]: #takes in a dataframe with an column containing the predicted wine quality
#and returns the error rate
def validation(df_after_prediction):
    error = 0
    num_row = 0
    # iterate through each row in df
    for idx, row in df_after_prediction.iterrows():
        num_row = num_row+1
        if not row['quality'] == row['predicted quality']:
            error = error+1
    errorPercent = error/num_row
```

```

    #print(errorPercent)
    return errorPercent

```

In [39]: *#takes in a dataframe
#and return a list of rows in the dataframe without the quality column*

```

def get_row_list(df):
    rows_list=[]
    for idx, row in df.iterrows():
        curr_row = [row['fixed.acidity'],row['volatile.acidity'],
                    row['citric.acid'],row['residual.sugar'],
                    row['chlorides'], row['free.sulfur.dioxide'],
                    row['total.sulfur.dioxide'], row['density'],
                    row['pH'], row['sulphates'],
                    row['alcohol']]
        rows_list.append(curr_row)
    return rows_list

```

In [40]: *#takes in: the training dataframe, the testing dataframe, the k value,
#the testing dataframe's row list, a kd-tree and returns the result of validation*

```

def getErrorPercent(df_training, df_testing, k_value, rows_list, tree):
    inds = tree.query(rows_list, k=k_value, return_distance = False)
    df_col_pred = predict_rate(df_training, inds)
    df_testing['predicted quality'] = df_col_pred
    return validation(df_testing)

```

We want to use the KD-tree in the sklearn library to predict the quality of a wine given its list of features(ingredients). A K-Dimensional tree is a tree data structure for organizing sample points in a k dimensional space. In our project, we have 11 dimensions as there are 11 features for each wine. A KDT will store our sample points (the rows of wine) according to all their features. Therefore, by querying the closest “neighbors” of a wine from the KDT, we get a list of other wines in the KDT positioned the closest to the query point. Since we have information about those wines’ quality, we can then predict the query wine’s quality by taking the quality with the highest occurrence among the nearest neighbors.

Now, we will build the KDT using the training dataset.

In [41]: *#list of rows of the training dataframe*

```

rows_list_training = get_row_list(df_training)

```

```

#the KD Tree that stores all the wines in the training set
tree = KDTree(rows_list_training)

```

After building the KD tree, we want to know how many numbers of neighbors we should query from it in order to get the most accurate prediction. Our initial guess is 5, but we want to actually go through more possible values to make sure we choose a best k value.

In [42]: *#going through possible k values [1,40]*

```

k_val = np.arange(1, 41, 1)
err_val=[]
for i in k_val:

```

```

df_training_cpy = df_training.copy()
err_val.append(getErrorPercent(df_training, df_training_cpy,
                                i ,rows_list_training, tree))

df_err = pd.DataFrame()
df_err['k value'] = k_val
df_err['error'] = err_val

```

In [43]: df_err[:10]

```

Out[43]:
   k value  error
0        1  0.000000
1        2  0.000000
2        3  0.000000
3        4  0.000000
4        5  0.000000
5        6  0.003518
6        7  0.013292
7        8  0.034011
8        9  0.061376
9       10  0.097733


```

```

In [44]: df_err
df_err.plot.scatter(x='k value', y='error')

```

Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x1a0d745320>



After checking through a set of k value candidates, we decide that k=5 would be the best choice. Because it gives us a very low error. Although k values less than 5 seem to have even lower error rates, there are a few reasons why we did not choose them:

- 1) We do not want to choose an even k value because we will then need to worry about having a tie.
- 2) k=1 and k=3 are odd numbers, but still not nice choices as they are too small.

We will now query 5 nearest neighbors for each wine instance in our training data set, predict the quality, and add to the dataframe.

```
In [45]: #the list of rows of the training dataset
rows_list_training = get_row_list(df_training)

#the list of lists of nearest neighbor indices
inds_k_5_training = tree.query(rows_list_training, k=5, return_distance = False)

#the list of predicted quality of each wine
df_training['predicted quality'] = predict_rate(df_training, inds_k_5_training)
```

/Users/aitingsieh/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: SettingWithCopyError: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
In [46]: df_training[:5]
```

```
Out[46]:
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	\
0	6.5	0.28	0.33	15.7	0.053	
1	8.3	0.31	0.39	2.4	0.078	
2	6.8	0.18	0.37	1.6	0.055	
3	6.4	0.30	0.33	5.2	0.050	
4	9.1	0.28	0.48	1.8	0.067	

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates	\
0	51.0	190.0	0.99780	3.22	0.51	
1	17.0	43.0	0.99444	3.31	0.77	
2	47.0	154.0	0.99340	3.08	0.45	
3	30.0	137.0	0.99304	3.26	0.58	
4	26.0	46.0	0.99670	3.32	1.04	

	alcohol	quality	predicted quality
0	9.7	6	6
1	12.5	7	7
2	9.1	5	5
3	11.1	5	5
4	10.6	6	6

We will compute the error percentage of our prediction.

```
In [47]: error_training = getErrorPercent(df_training, df_training, 5,
                                         rows_list_training, tree)
         error_training
```

```
/Users/aitinghsieh/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:6: SettingWithC
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out[47]: 0.0
```

Now that our model is ready, we will go ahead and start predicting the quality of wines in the testing dataset.

```
In [48]: #the list of rows in the testing set
         rows_list_testing = get_row_list(df_testing)

         #the list of lists of nearest neighbor indices
         inds_k_5_testing = tree.query(rows_list_testing, k=5, return_distance = False)

         #the list of predicted quality of each wine
         df_testing['predicted quality'] = predict_rate(df_training, inds_k_5_testing)
```

```
/Users/aitinghsieh/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: SettingWithC
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
In [49]: # err_testing = getErrorPercent(df_training, df_testing, 5, rows_list_testing)
         err_testing = validation(df_testing)
         err_testing
```

```
Out[49]: 0.4609375
```

7 Part 4 - Data Analysis / Results

It turned out that the error percentage is pretty big, opposite from what we were expecting. Therefore, now we want to see how accurate our prediction is.

```
In [50]: df_testing[:10]
```

```

Out [50]:      fixed.acidity  volatile.acidity  citric.acid  residual.sugar  chlorides  \
2558           7.9           0.54           0.34           2.50           0.076
2559           7.3           0.30           0.34           2.70           0.044
2560           6.7           0.46           0.24           1.70           0.077
2561           6.7           0.26           0.26           4.10           0.073
2562           7.4           0.19           0.30           1.40           0.057
2563           5.7           0.22           0.22          16.65           0.044
2564          12.3           0.27           0.49           3.10           0.079
2565           7.3           0.91           0.10           1.80           0.074
2566           9.3           0.40           0.49           2.50           0.085
2567           6.1           0.20           0.40           1.90           0.028

      free.sulfur.dioxide  total.sulfur.dioxide  density    pH  sulphates  \
2558                8.0                17.0  0.99235  3.20        0.72
2559               34.0               108.0  0.99105  3.36        0.53
2560               18.0                34.0  0.99480  3.39        0.60
2561               36.0               202.0  0.99560  3.30        0.67
2562               33.0               135.0  0.99300  3.12        0.50
2563               39.0               110.0  0.99855  3.24        0.48
2564               28.0                46.0  0.99930  3.20        0.80
2565               20.0                56.0  0.99672  3.35        0.56
2566               38.0               142.0  0.99780  3.22        0.55
2567               32.0               138.0  0.99140  3.26        0.72

      alcohol  quality  predicted quality
2558     13.1       8           6
2559     12.8       8           6
2560     10.6       6           6
2561      9.5       5           4
2562      9.6       6           7
2563      9.0       6           6
2564     10.2       6           4
2565      9.2       5           5
2566      9.4       5           5
2567     11.7       5           6

```

It looks like some wines have the same predicted and actual quality while some others don't. We will need to do some more work to see how our algorithm performs.

```

In [51]: import matplotlib.pyplot as plt

fig = plt.figure()
fig.set_size_inches(18.5, 10.5)
ax1 = fig.add_subplot(111)

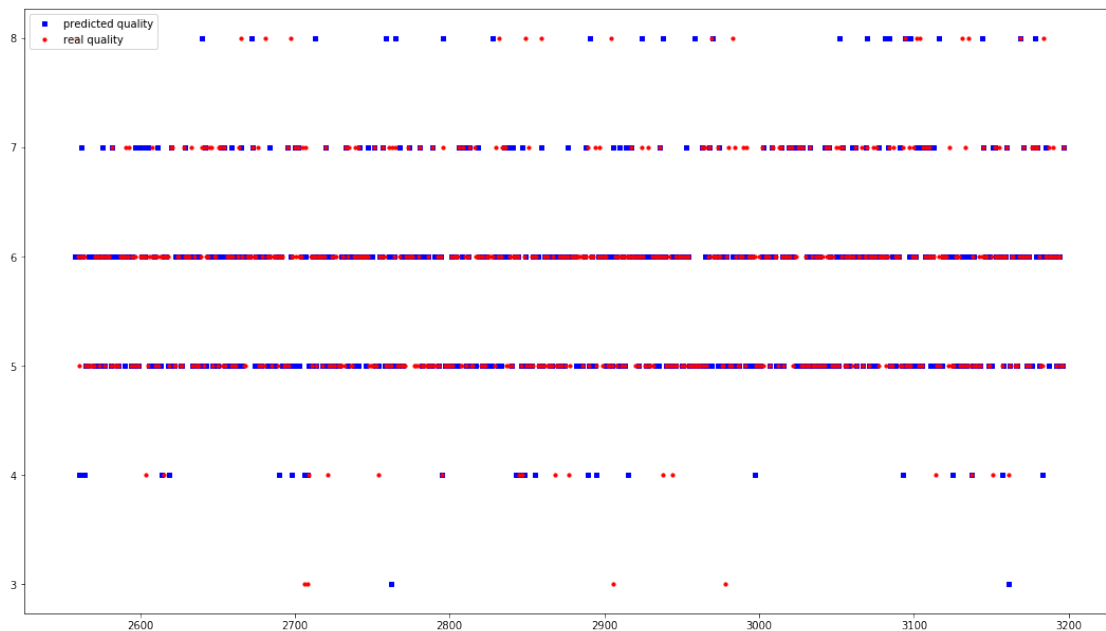
ax1.scatter(df_testing.index, df_testing['predicted quality'],
            s=10, c='b', marker="s", label='predicted quality')
ax1.scatter(df_testing.index, df_testing['quality'], s=10, c='r',

```

```

        marker="o", label='real quality')
plt.legend(loc='upper left')
plt.show()

```



The scatter plot above does not look so nice and we cannot really tell much about our prediction by looking at it. Thus, we will calculate the root mean square error we got and store the value in our dataframe.

In [52]: *#takes in two values and calculautes their root mean square difference
#and then returns that result*

```

def rmse(pred, tar):
    diff = pred - tar
    diff_sq = diff ** 2
    mean_diff_sq = diff_sq.mean()
    rmse_val = np.sqrt(mean_diff_sq)
    return rmse_val*1.0

```

```

In [53]: df_testing['rms']=0.0
        for i in df_testing.index:
            df_testing.at[i,'rms'] = rmse(df_testing.at[i,'predicted quality'],
                                           df_testing.at[i,'quality'])

```

/Users/aitingsieh/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCovarianceWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
 """Entry point for launching an IPython kernel.

```
In [54]: df_testing[:10]
```

```
Out[54]:
```


	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	\
2558	7.9	0.54	0.34	2.50	0.076	
2559	7.3	0.30	0.34	2.70	0.044	
2560	6.7	0.46	0.24	1.70	0.077	
2561	6.7	0.26	0.26	4.10	0.073	
2562	7.4	0.19	0.30	1.40	0.057	
2563	5.7	0.22	0.22	16.65	0.044	
2564	12.3	0.27	0.49	3.10	0.079	
2565	7.3	0.91	0.10	1.80	0.074	
2566	9.3	0.40	0.49	2.50	0.085	
2567	6.1	0.20	0.40	1.90	0.028	

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates	\
2558	8.0	17.0	0.99235	3.20	0.72	
2559	34.0	108.0	0.99105	3.36	0.53	
2560	18.0	34.0	0.99480	3.39	0.60	
2561	36.0	202.0	0.99560	3.30	0.67	
2562	33.0	135.0	0.99300	3.12	0.50	
2563	39.0	110.0	0.99855	3.24	0.48	
2564	28.0	46.0	0.99930	3.20	0.80	
2565	20.0	56.0	0.99672	3.35	0.56	
2566	38.0	142.0	0.99780	3.22	0.55	
2567	32.0	138.0	0.99140	3.26	0.72	

	alcohol	quality	predicted	quality	rms
2558	13.1	8	6	2.0	
2559	12.8	8	6	2.0	
2560	10.6	6	6	0.0	
2561	9.5	5	4	1.0	
2562	9.6	6	7	1.0	
2563	9.0	6	6	0.0	
2564	10.2	6	4	2.0	
2565	9.2	5	5	0.0	
2566	9.4	5	5	0.0	
2567	11.7	5	6	1.0	

```
In [55]: df_testing['rms'].plot.hist()
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x1a0d91fc18>
```

We can see that the error we get is actually very small. The predicted quality is in the most case actual quality ≤ 1 , with a few ones having a more inaccurate result. So after all, our prediction is still not bad.

8 Part 5 - Ethics and Privacy

In our ventures in trying to determine what predicts wine quality, it is important to consider who and what may be negatively impacted by the results we obtain. Our results may reveal that certain qualities in wine are more favorable than others, affecting the business of companies that may be known to have more or less of a favorable quality. For example, if it is found that a wine with a higher sugar content is less favorable, companies and vineyards known for sweeter wines may find their businesses negatively impacted. Additionally, if a certain company claims the key to quality is one specific physicochemical property and we find their information to be false, that company may also find their business to be negatively impacted. With more awareness as to what makes a quality wine, perceived high quality wine price could drop and perceived low quality price could increase, if our results dispute their perception.

Additionally, it is important to consider that our data is limited, biased, and skewed. We cannot be sure that the data was collected in a way that is accurate, precise, and equitable. Since all our data is from one wine company, we are sure that all variations of wine in the world and that all physicochemical properties affecting taste are not represented in our dataset. With this in mind, if our results grow in popularity, it can lead to positive or negative business for the company who provided this data, depending on results. Lastly, we know that wine quality is subjective and the "rating" portion of our data is biased. With biased ratings, it is also likely that the results are skewed, with more ratings at extreme ends than close to a median. With all this in mind, we also consider that our results cannot be fully used to defend a hypothesis, and that outside evidence is necessary for full confirmation.

9 Part 6 - Conclusion & Discussion

In conclusion, we found that though we are able to somewhat predict wine quality, it is still not fine tuned enough to use ethically. Our prediction is generally on par, sometimes slightly off, than the real quality, within 2 at most. In relation to our hypothesis, all our factors of physicochemical features combine to create different optimal wines, meaning that all of our factors are involved in creating the proper flavor profile. Our original hypothesis predicts that lower volatile acidity, medium alcohol percentage, and higher citric acid would create the best quality, and we did not necessarily find that to be true. Interestingly, though we are already working with 12 variables, there are still contributors to wine quality that are missing. With each of the 12, we were unable to find one that stood out or a few that stood out together in making a quality wine. They all needed to combine with one another in a grand way to make a high quality wine.

A possible confound is in variables that cannot directly be measured from the wine itself. Many experts cite the climate and soil at the vineyard in which the wine originated, something that was not taken into account with our data. This may account for part of the inaccuracy we faced. Limitations in our data may also account for inaccuracy. As stated before, we cannot be sure of the precision with which the data was collected. Additionally, all of the wine in our dataset came from the same vineyard, limiting our scope. Our dataset also lacked wine at the extreme ends of high and low quality, limiting our training, thus leading to inaccuracy. A few confounds are inevitable, but it may have done us well to try to alleviate some of their effects.

If we were to continue this project, it may help to look deeper in to advanced data science techniques for tackling problems like this. In our pursuit to predict what makes a high quality wine, we found and learned about KD Tree, which made sense of our data when we could not. There may be more techniques out there that would've better our predictions and help us properly pinpoint the best wine. Additionally, the addition of more data may have helped us better find the correlations we were looking for. Despite our attempts, the one true perfect wine is still somewhat of a mystery, but we are sure data scientists will tackle it better than we could.